

Assignment 1

Decision Trees:

- Classify by walking tree
- Greedy recursive splitting
- Information gain as score, entropy
- $O(m \log n)$,

Decision Stump Rule Search:

- Ignore rules outside feature ranges.
- Sort examples $O(n \log n)$
- $O(n)$ score updates
- Total cost: $O(nd \log n)$, data size: $O(nd)$

Learning Theory:

- THE TEST DATA CANNOT INFLUENCE THE TRAINING PHASE IN ANY WAY.
- Fundamental Trade-Off:
 1. How small you can make the training error.
 2. How well training error approximates the test error.
- Simple models have good approximation of test error, but fit training data poorly.
- Cross validation, test on your training data, use each point as test once.
- No free lunch theorem: no best machine learning model for every problem

Naive Bayes:

$$p(y_i = \text{"spam"} | x_i) = \frac{p(x_i | y_i = \text{"spam"}) p(y_i = \text{"spam"})}{p(x_i)}$$

- Compute $p(y = \text{"spam"} | x_i) > p(y = \text{"not spam"} | x_i)$
- Need to estimate $p(x_i | y = \text{"spam"})$
- Runtime: Test: $O(kd)$; Train: $O(nkd + n)$

Bayes Rule:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Descriptive Statistics:

- Mean, Median, Quantiles (v s.t. $x\% < v$)
- Box Plot

Assignment 2

KNN:

- Non-parametric
- Distance fn: Euclidean, L1 Dist, Jaccard similarity, Cosine, Distance after dimension reduction, Metric learning

$$\text{Jaccard similarity: } D(x_1, x_2) = (x_1 \cup x_2) / (x_1 \cup x_2)$$

Ensemble Methods: Boosting, Averaging

Random Forests:

- Train number of deep decision trees on bootstrapped samples from the main dataset + small random subset of features on each stump
- Very fast, good out of the box classifier
- Bayesian Model Averaging
- Train $O(tm \log n)$, $t = \# \text{trees}$, $m = \text{depth}$, $d = \# \text{features}$, can be \sqrt{d}
- Classify one is $O(tm)$

Clustering: No best method, no test error

K-Means:

- k clusters, random center
- add one point and update mean based on assignment
- initialization important, need to know k
- Total cost: $O(ndk)$, updating means $O(nd)$
- vector quantization: cluster colors
- K-Medians: Can also use L1-norm + median
- K-Means++, init clusters by furthest distance
- Convex clusters

Density based clustering:

- Non-convex clusters

DBSCAN:

- Radius, if at least minPoints of same type become "core"
- Merge clusters if reachable
- Can have outliers
- Sensitive to boundaries
- need lots of points in high dimensions
- finding cluster is expensive
- $O(n^2d)$

Lp-norms: L1, L2, L-inf

$$\|x\|_p := \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}.$$

Elbow method: choose k in which the the sharpest elbow (biggest change in slope) for minimum error vs k.

Assignment 3

Linear regression (adding bias, change of basis, RBFs, regularization, cost of training/testing, effect of basis parameters).

- 1D $f(w) = \sum_{i=1}^n (wx_i - y_i)^2$
- Minimize least squares $w = \frac{\sum_{i=1}^n y_i x_i}{\sum_{i=1}^n x_i^2}$
- Solution
- Bias variable, add row with 1s
- Change of basis, linear least squares
- each row: $[1 \times x^2 \dots x^p]; [1 \times \sin(6x)]$
- Radial basis functions: non-parametric

$$\begin{bmatrix} g(\|x_1 - x_1\|) & g(\|x_1 - x_2\|) & \dots & g(\|x_1 - x_n\|) \\ g(\|x_2 - x_1\|) & g(\|x_2 - x_2\|) & \dots & g(\|x_2 - x_n\|) \\ \vdots & \vdots & \ddots & \vdots \\ g(\|x_n - x_1\|) & g(\|x_n - x_2\|) & \dots & g(\|x_n - x_n\|) \end{bmatrix}$$

$$g(\alpha) = \exp\left(-\frac{\alpha^2}{2\sigma^2}\right)$$

- Can combine both

$$f(w) = -\log(p(w|X, y)) = -\sum_{i=1}^n \log(p(y_i|x_i, w)) - \sum_{j=1}^d \log(p(w_j)) + \text{const.}$$

Robust regression (weighted least squares, smooth approximations, computing gradients).

- Robust means fine with extreme outliers



- Absolute is most robust convex error function
- Weighted least squares: weight on each training example

Convexity:

- 1-variable, twice-differentiable function is convex iff $f''(w) \geq 0$ for all 'w'.
- A convex function multiplied by non-negative constant is convex.
- Norms and squared norms are convex.
- The sum of convex functions is a convex function.
- The max of convex functions is a convex function.
- Composition of a convex function and a linear function is convex.
- Not true: composition of convex and convex is convex

0-1 Loss: Number of classification errors

- Hinge: Convex, $\max\{0, 1 - y_i w^T x_i\}$
- SVM is hinge loss + L2-regularizer
- Logistic: smooth approximation to hinge, differentiable

Assignment 4:

Convex and MLE/MAP estimation (showing functions are convex, connection between probabilities and losses/regularizers)

- likelihood function: $p(y|X, w)$
- minimize negative log-likelihood

$$f(w) = -\sum_{i=1}^n \log(p(y_i|x_i, w))$$

Maximum a Posteriori (MAP) Estimation:

$$p(w|X, y) \propto p(y|X, w)p(w)$$

- $p(w|X, y)$ "posterior"
- $p(y|X, w)$ "likelihood"
- $p(w)$ "prior"

- Loss = NLL
- Regularizer = negative log-prior

Sigmoid: $h(z) = 1/(1+\exp(-z))$

Logistic Regression:

$$p(y_i|x_i, w) = \frac{1}{1+\exp(-y_i w^T x_i)}$$

$$f(w) = \sum_{i=1}^n \log(1 + \exp(-y_i w^T x_i)) + \frac{\lambda}{2} \|w\|^2$$

Regularizers:

- "robust" = less influenced by large values
- L0 sparsity, non-convex
- L1 makes things more sparse, non-convex
- Huber loss:
$$L_\delta(a) = \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta, \\ \delta(|a| - \frac{1}{2}\delta), & \text{otherwise.} \end{cases}$$
- L2
- L Infinite = max, log-sum-exp is smooth approx
- LSE(x_1, \dots, x_n) = $\log(\exp(x_1) + \dots + \exp(x_n))$
- Frobenius Norm: sqrt(each element squared)

Regularized logistic regression (loss functions for binary classification, effect of different regularizers on overfitting and sparsity)

Vectors/matrices/norms (summation and vector/matrix/norm notation, minimizing quadratic functions as linear systems)

Multi-class Logistic (one-vs-all, softmax loss derivatives and implementation)

- multi-class y_i is row of d [+1/-1]
- multi-class classification, exactly one +1, encode as class k
- can take max of y_i to get correct classification

softmax/multinomial logistic regression

$$p(y_i|x_i, w) = \frac{\exp(w_{y_i}^T x_i)}{\sum_{c=1}^k \exp(w_c^T x_i)}$$

- logistic regression for binary labels
- softmax for multi-class
- ordinal logistic regression, thresholds of sigmoid as params

Assignment 5:

Principle Component Analysis

PCA (computing 1st PC, scaling issue, PCA for visualization, PCA for compression)

- $\|w\|_2 = 1$
- $X = ZW$
- $f(Z, W) = \sum_{i=1}^n (w^T z_i - x_i)^2$
- reduce dimensionality

$$f(W, Z) = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^d (w_j^T z_i - x_{ij})^2 = \frac{1}{2} \|ZW - X\|_F^2$$

- Solve via:
 - Singular value decomposition (SVD) non-iterative
 - Alternate between updating W and Z
 - Stochastic gradient / Gradient Descent
- X : n by d , Z : n by k , W : k by d

Beyond PCA (different loss functions, effect of regularizers)

- PCA minimized approximation error
- PCA maximizes variance
- $var(X) = \frac{1}{n} \|X\|_F^2$
- Regularizers can split into different components, sparsity
- Non-negative matrix factorization
 - Regularizes, negative can't cancel out large values

Multi-dimensional scaling (basic model, ISOMAP and geodesic distance)

- Directly optimize the location of z_i values
- Gradient descent
- non-convex, sensitive to initialization

$$f(Z) = \sum_{i=1}^n \sum_{j=i+1}^n d_3(d_2(z_i, z_j) - d_1(x_i, x_j))$$

ISOMAP: Geodesic distance

- Use neighbors to compute edge weights
- Use dijkstra's to find shortest path distances

t-SNE: wayyy better, but ugly hack

- focus on small distances
- allow large variance in large distances
- not always better than PCA, can "twist" the plane

Assignment 6:

PageRank (random walk model)

- Probability of landing on page as $t \rightarrow \infty$
- Add small probability of going to random webpage
- At each step follow random link on page

Stochastic gradient and neural networks

- effect of step-size: too large = can't converge, too small = can't get to the solution
- effect of standardization: helps typically
- effect of initialization: random weights work better with non-convex losses
- effect of non-linearity: makes it a universal approximator
- effect of regularization: higher training error, better estimate of test error
- Dropout: randomly set some $x_i, z_i = 0$
- Convolutional
- Vanishing gradient, ReLU (hinge loss)

Inner product

$$a^T b = b^T a$$

$$a^T (b + c) = a^T b + a^T c$$

$$\text{Orthogonal } a^T b = 0 \quad \text{Orthonormal } a^T b = 1$$

Matrix Multiplication

$$A(BC) = (AB)C$$

$$A(B + C) = AB + AC$$

$$AB \neq BC$$

$$(AB)^T = B^T A^T$$

$$(AB)^2 = ABAB$$

$$\text{Matrix} \times \text{vector} = \text{vector}$$

$$x^T A y = x^T (A y) = (A y)^T x = y^T A^T x$$

Inverses

$$A^{-1} A = I = A A^{-1}$$

$$(A^{-1})^T = (A^T)^{-1}$$

$$(\gamma A)^{-1} = \gamma^{-1} A^{-1}$$

$$(AB)^{-1} = B^{-1} A^{-1}$$

Assignment 1:

Summary statistics (like range, median, and quantiles).

Data visualization (histogram, scatterplot, and box plot).

Decision trees (how to fit stumps using classification error, how to classify a new example, runtime in $O()$ notation, effect of depth).

Learning theory (training vs. test error, validation sets and cross-validation, fundamental trade-off)

Naive Bayes (conditional probability, fitting the probabilities, classifying a new example, runtime).

Assignment 2:

K-nearest neighbours (how to classify a new example, runtime in $O()$ notation, effect of k , condensed version).

Random forests (what bootstrap does, how random trees works, effect of number of trees/random-features, how to classify new example).

K-Means (effect of initialization, error functions, runtime, how to cluster new example, elbow method, k-medians, vector quantization).

Density-based clustering (effect of parameters, shape of clusters).

Assignment 3:

Vectors/matrices/norms (summation and vector/matrix/norm notation, minimizing quadratic functions as linear systems)

Linear regression (adding bias, change of basis, RBFs, regularization, cost of training/testing, effect of basis parameters).

Robust regression (weighted least squares, smooth approximations, computing gradients).

Assignment 4:

Regularized logistic regression (loss functions for binary classification, effect of different regularizers on overfitting and sparsity)

Convex and MLE/MAP estimation (showing functions are convex, connection between probabilities and losses/regularizers)

Multi-class Logistic (one-vs-all, softmax loss derivatives and implementation)

Assignment 5:

PCA (computing 1st PC, scaling issue, PCA for visualization, PCA for compression)

Beyond PCA (different loss functions, effect of regularizers)

Multi-dimensional scaling (basic model, ISOMAP and geodesic distance)

Assignment 6:

PageRank (random walk model)

Stochastic gradient and neural networks (effect of step-size, effect of standardization, effect of initialization, effect of non-linearity, effect of regularization)

Operation	Input	Output	Algorithm	Complexity
Matrix multiplication	Two $n \times n$ matrices	One $n \times n$ matrix	Schoolbook matrix multiplication	$O(n^3)$
			Strassen algorithm	$O(n^{2.807})$
			Coppersmith–Winograd algorithm	$O(n^{2.376})$
			Optimized CW-like algorithms ^{[14][15][16]}	$O(n^{2.373})$
Matrix multiplication	One $n \times m$ matrix & one $m \times p$ matrix	One $n \times p$ matrix	Schoolbook matrix multiplication	$O(nmp)$
Matrix inversion*	One $n \times n$ matrix	One $n \times n$ matrix	Gauss–Jordan elimination	$O(n^3)$
			Strassen algorithm	$O(n^{2.807})$
			Coppersmith–Winograd algorithm	$O(n^{2.376})$
			Optimized CW-like algorithms	$O(n^{2.373})$
Singular value decomposition	One $m \times n$ matrix	One $m \times m$ matrix, one $m \times n$ matrix, & one $n \times n$ matrix		$O(mn^2)$ ($m \leq n$)
		One $m \times r$ matrix, one $r \times r$ matrix, & one $n \times r$ matrix		
Determinant	One $n \times n$ matrix	One number	Laplace expansion	$O(n!)$
			Division-free algorithm ^[17]	$O(n^4)$
			LU decomposition	$O(n^3)$
			Bareiss algorithm	$O(n^3)$
			Fast matrix multiplication ^[18]	$O(n^{2.373})$
Back substitution	Triangular matrix	n solutions	Back substitution ^[19]	$O(n^2)$